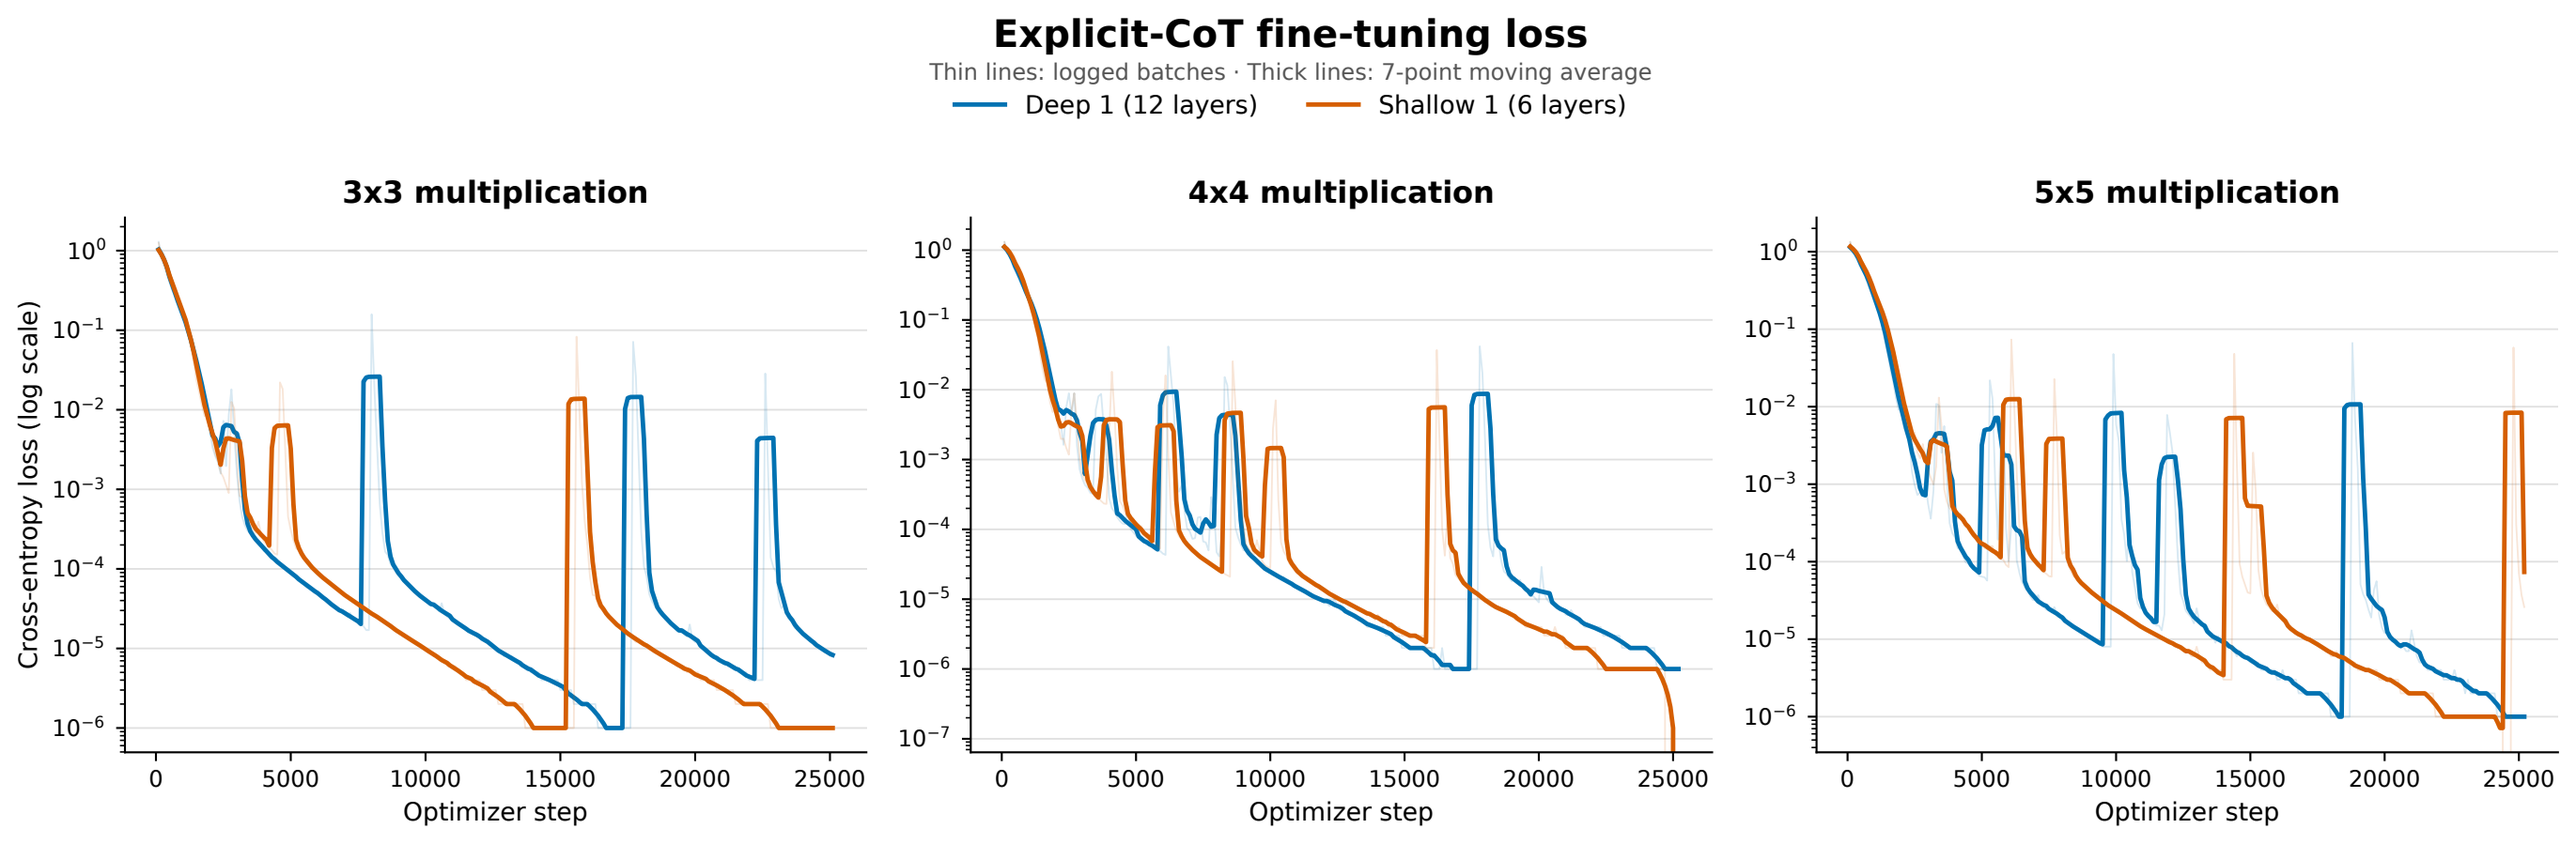
Explicit-CoT fine-tuning loss
Thin lines: logged batches · Thick lines: 7-point moving average
Deep 1 (12 layers)
Shallow 1 (6 layers)
3x3 multiplication
4x4 multiplication
5x5 multiplication
Cross-entropy loss (log scale)
Optimizer step
0
5000
10000
15000
20000
25000